

Micrium

© Copyright 2010, Micrium
All Rights reserved

**New Features and Services
since**

μC/OS-II V2.00

(Current Version: V2.92)

www.Micrium.com

Introduction

This document describes all the features and services added to **μC/OS-II** since the introduction of the hard cover book *MicroC/OS-II, The Real-Time Kernel*, ISBN 0-87930-543-6. The software provided with the book was version 2.00 or V2.04. The version number of the change is shown when appropriate.

Delete task on incorrect return (V2.89)

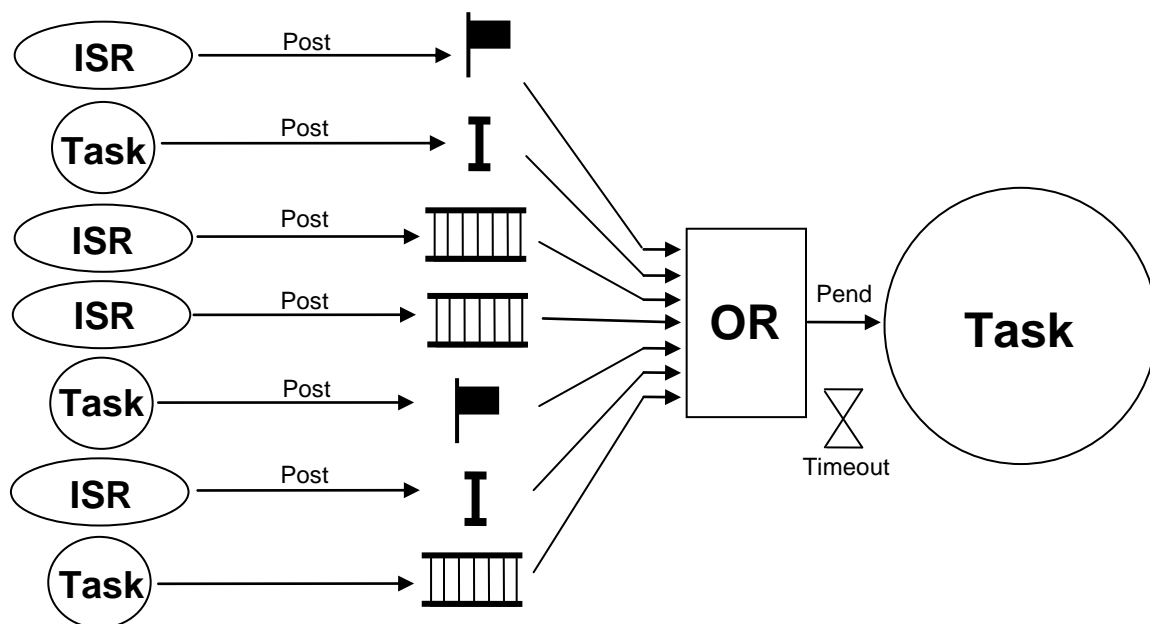
μC/OS-II now contains a new function called `OS_TaskReturn()`. All **μC/OS-II** tasks are not allowed to return. If a task returns by mistake, `OS_TaskReturn()` catches those and deletes the task.

`OS_TaskReturn()` calls `OSTaskReturnHook()` which in turn calls `App_TaskReturnHook()`.

Pend on Multiple Events (V2.86)

μC/OS-II now contains a new function called `OSEventPendMulti()` which allows a task to pend on multiple events (semaphores, mailboxes and queues) in any combination (see example diagram below). This new function is found in `OS_CORE.C` and is enabled by setting `OS_EVENT_MULTI_EN` to 1 in `OS_CFG.H`.

With `OSEventPendMulti()` it's possible to pend on any number of semaphores, mailboxes and message queues at the same time (we don't support Mutex and Event Flags at this time). If a task pends on a combination of the above 'events' then, as soon an event is posted (and the pending task is the highest priority task pending on the event), the waiting task will wake up and be 'handed' the event. If events are present as the task pends then ALL the available events will be provided to the task.



Timer Manager (V2.81)

μC/OS-II now provides support for periodic as well as one-shot timers. This functionality is found in `OS_TMR.C` and is enabled by setting `OS_TMR_EN` to 1 in `OS_CFG.H`. Your application can have any number of timers (up to 65500). When a timer times out, an optional callback function can be called allowing you to perform any action (signal a task, turn on/off a light, etc.). Each timer has its own callback function.

IMPORTANT

The APIs for the Timer Manager were changed in V2.83 from what they were in V2.81 and V2.82. This was necessary to correct some issues with the Timer Manager. Please consult the Reference Manual for the new APIs.

When timer management is enabled, **μC/OS-II** creates a timer task (`OSTmrTask()`) which is responsible for updating all the timers. The priority of this task is determined by `OS_TASK_TMR_PRIO` which should be defined in your application's `APP_CFG.H`.

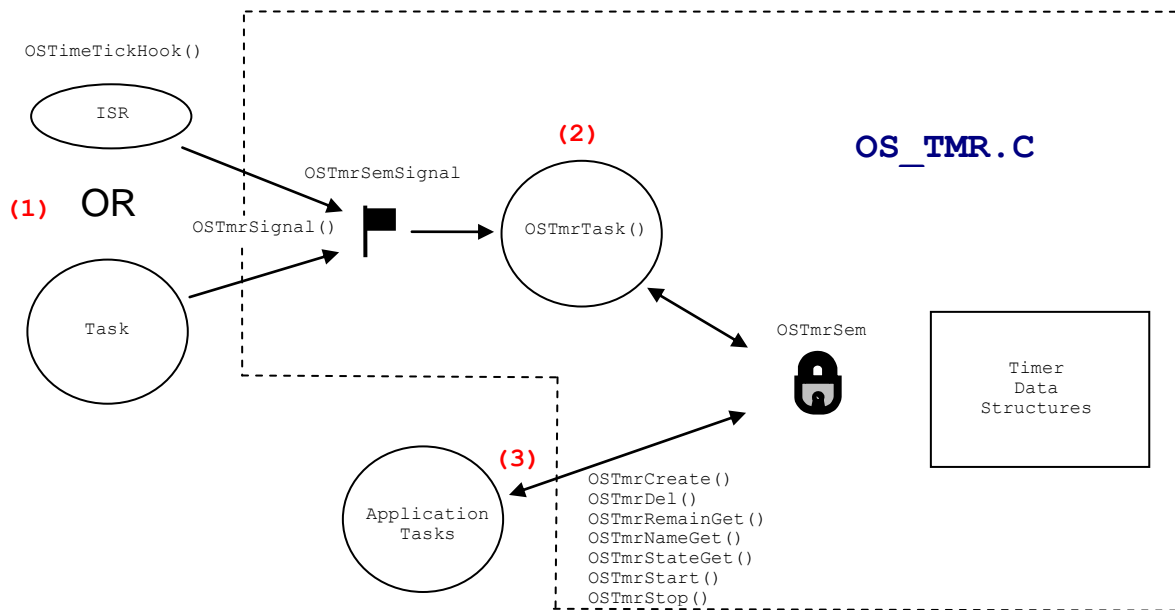
The timer manager provides a number of services to your applications. Specifically, you can call one of the following functions (see the **μC/OS-II** reference manual for a description of these functions) from your tasks:

| | |
|-------------------------------|--|
| <code>OSTmrCreate()</code> | Create a timer |
| <code>OSTmrDel()</code> | Delete a timer |
| <code>OSTmrRemainGet()</code> | Determine how much time before a timer expires |
| <code>OSTmrNameGet()</code> | Get the name of a timer |
| <code>OSTmrStateGet()</code> | Get the state of a timer (UNUSED, STOPPED, RUNNING, COMPLETED) |
| <code>OSTmrStart()</code> | Start a timer |
| <code>OSTmrStop()</code> | Stop a timer |

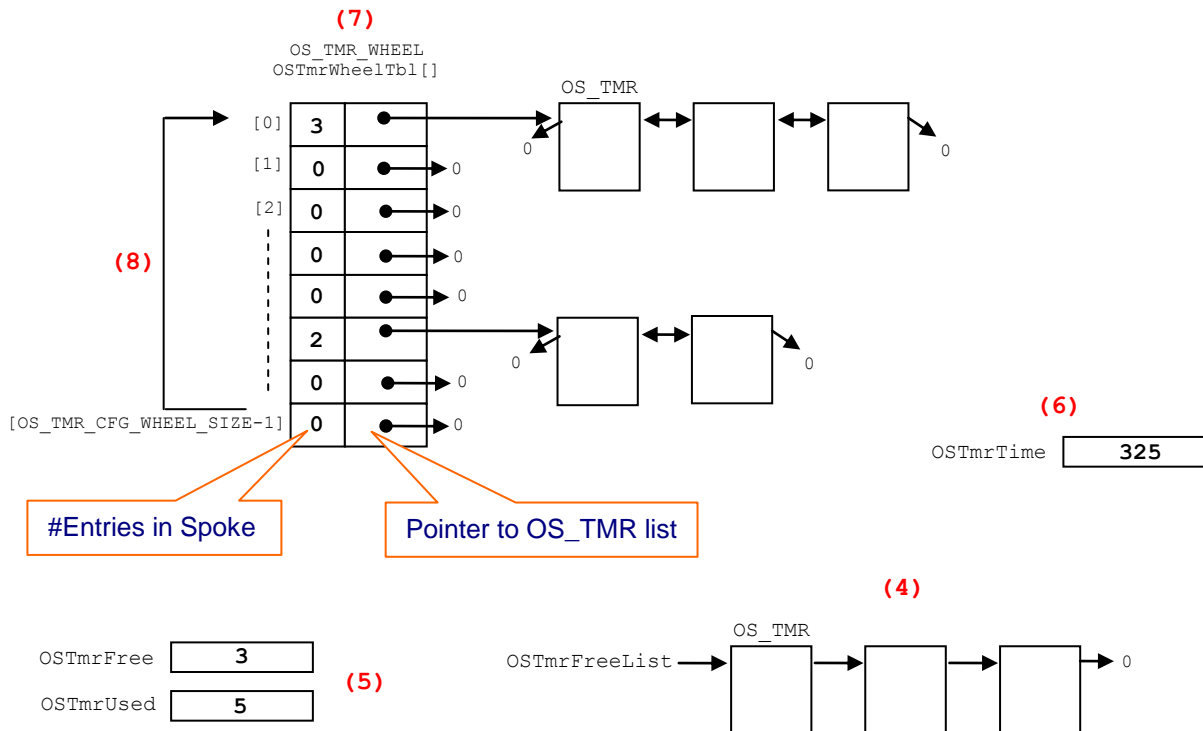
You should note that you **CANNOT** call these functions from ISRs.

The drawing below shows the task model of the Timer Manager. You should note that semaphore management needs to be enabled (you need to set `OS_SEM_EN` to 1 in `OS_CFG.H`) for the timer manager to work. The timer manager requires two (2) semaphores.

- (1) An ISR or an application task needs to ‘signal’ a counting semaphore by calling `OSTmrSignal()` at the rate specified in `OS_TMR_CFG_TICK_RATE` (see `OS_CFG.H`). The counting semaphore is called `OSTmrSemSignal` that is initialized to 0 by **µC/OS-II** when `OSInit()` is called. You should note that you should **ONLY** call `OSTmrSignal()` and not worry about the semaphore; it’s encapsulated by `OSTmrSignal()`.
- (2) The timer management task (`OSTmrTask()`) pends forever on the counting semaphore waiting for it to be signaled. When the semaphore is signaled, `OSTmrTask()` acquires another semaphore (a binary semaphore in this case, `OSTmrSem`) to gain exclusive access to timer data structures. When `OSTmrTask()` is the owner of the semaphore it updates all the timers created by your application.
- (3) Your application accesses timer data structure via interface functions. These functions allow you to create, delete, start and stop timers as well as examine the amount of time remaining before a timer times out.



The drawing below shows the data structures used in the timer manager.



- (4) Each timer is characterized by a data structure of type **OS_TMR** (see **ucos_ii.h**). Each timer contains the 'period' of the timer (if the timer is to operate in periodic mode), the name of the timer, a timer 'match' value (described later) and other fields used to link the timer. Free timers are placed in a singly linked list of 'unused' timers pointed to by **OSTmrFreeList**.

- (5) The number of free timers is held in `OSTmrFree` and the number of used (or allocated) timers is held in `OSTmrUsed`. Of course, the total number of timers is the sum of these two fields and, unless you don't properly use the timer management services, that sum should always equal `OS_TMR_CFG_MAX`.
- (6) Every time `OSTmrSignal()` is called, the unsigned 32-bit variable `OSTmrTime` is incremented by one and used to see if timers have expired.
- (7) The timer manager keeps track of which timer it needs to update using a 'timer wheel'. The wheel is basically an array of structures of type `OS_TMR_WHEEL` (see `ucos_ii.h`) that wraps around. This structure contains two fields: a pointer to a doubly-linked list of `OS_TMR` structures and, the number of entries in that list.
- (8) The 'wheel' contains `OS_TMR_CFG_WHEEL_SIZE` entries or spokes.

`OS_TMR` structures are inserted in the wheel when you call `OSTmrStart()`. The position (i.e. spoke) in `OSTmrWheelTbl[]` for a specific timer is given by:

```
match = OSTmrTime + period;
spoke = match % OS_TMR_CFG_WHEEL_SIZE;
```

The 'match' corresponds to the value that `OSTmrTime` needs to reach before the timer expires. For example, let's say that `OSTmrTime` is 0 (just initialized) and we want to create a timer that will expire every second (assuming `OS_TMR_CFG_TICKS_PER_SEC` is set to 10). Also, let's assume that `OS_TMR_CFG_WHEEL_SIZE` is 8 (as shown in the diagram above).

```
match = OSTmrTime + period;
match = 0 + 10;
match = 10;

spoke = match % OS_TMR_CFG_WHEEL_SIZE;
spoke = 10 % 8;
spoke = 2;
```

This means that `OSTmrStart()` will obtain a free `OS_TMR` data structure from the free list of timers and place this data structure in `OSTmrWheelTbl[]` at position #2 in the table. `OSTmrStart()` will then store the 'match' value in the `OS_TMR` data structure.

Every time `OSTmrTime` is incremented by `OSTmrTask()`, `OSTmrTask()` goes through ALL the `OS_TMR` structures placed at spoke (`OSTmrTime % OS_TMR_CFG_WHEEL_SIZE`) to see if `OSTmrTime` 'matches' the value stored in the `OS_TMR` structure. If a match occurs, the timer is removed from the list. If the timer was started by `OSTmrStart()` with a 'periodic' option then, the `OS_TMR` structure is placed in the `OSTmrWheelTbl[]` by calculating its new position, again using `OSTmrTime + period`. In our example, the new 'spoke' would be:

```

match = OSTmrTime + period;
match = 10 + 10;
match = 20;

spoke = match % OS_TMR_CFG_WHEEL_SIZE;
spoke = 20 % 8;
spoke = 4;

```

The use of a timer wheel basically reduces the execution time of the timer task so that it only handles a few of the timers. Of course, the worst case is such that all timers are placed in the same spoke of the timer wheel. However, statistically, this will occur rarely. It's generally recommended to keep the size of the wheel a fraction of the total number of times. In other words, you should set:

$$\text{OS_TMR_CFG_WHEEL_SIZE} \leq \text{Fraction of (OS_TMR_CFG_MAX)}$$

A fraction of 2 to 8 should work well.

RAM usage (in bytes) for the timer manager is shown below:

```

2 * sizeof(INT16U) +
1 * sizeof(INT32U) +
3 * sizeof(POINTER) +
OS_TASK_TMR_STK_SIZE * sizeof(OS_STK) +
OS_TMR_CFG_WHEEL_SIZE * (sizeof(INT16U) + sizeof(POINTER)) +
OS_TMR_CFG_MAX * (4 * sizeof(POINTER) +
                  2 * sizeof(INT32U) +
                  3 * sizeof(INT8U) +
                  OS_TMR_CFG_NAME_SIZE * sizeof(INT8U))

```

Because INT8Us and BOOLEANS are typically 1 byte, INT16Us are 2 bytes and INT32Us are 4 bytes, we can simplify the above equation as follows:

```

2 * 2 +
1 * 4 +
3 * sizeof(POINTER) +
OS_TASK_TMR_STK_SIZE * sizeof(OS_STK) +
OS_TMR_CFG_WHEEL_SIZE * (2 + sizeof(POINTER)) +
OS_TMR_CFG_MAX * (4 * sizeof(POINTER) +
                  2 * 4 +
                  3 +
                  OS_TMR_CFG_NAME_SIZE)

```

Or,

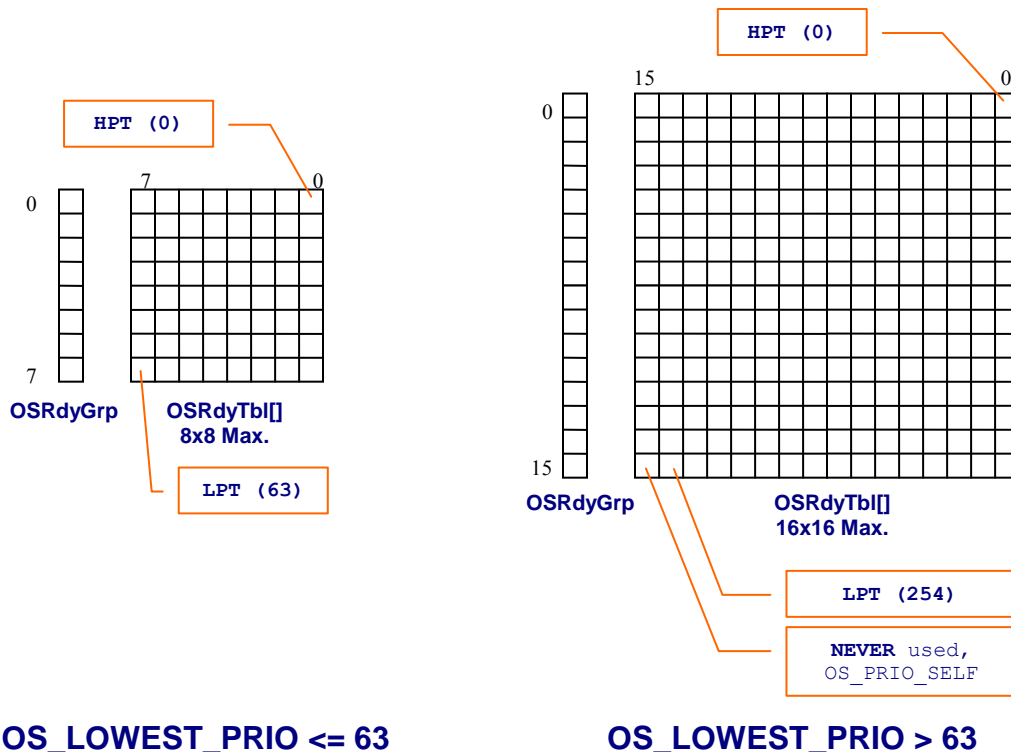
```

8 +
3 * sizeof(POINTER) +
OS_TASK_TMR_STK_SIZE * sizeof(OS_STK) +
OS_TMR_CFG_WHEEL_SIZE * (2 + sizeof(POINTER)) +
OS_TMR_CFG_MAX * (4 * sizeof(POINTER) + 11 + OS_TMR_CFG_NAME_SIZE)

```


Support for 255 tasks (V2.80)

μC/OS-II can now support up to **255** tasks. To support up to 255 tasks, we simply increased the ready list and event wait lists to a matrix of 16x16 instead of 8x8. In fact, the actual size of the matrix (whether 8x8 or 16x16) depends on the value of `OS_LOWEST_PRIO` in `OS_CFG.H`. If `OS_LOWEST_PRIO` is less than or equal to 63, we use an 8x8 matrix and thus **μC/OS-II** behaves exactly the same as it used to. If you specify a value for `OS_LOWEST_PRIO` to be greater than 63, we use the 16x16 matrix as show below.



You should note that the actual size of the matrix depends on `OS_LOWEST_PRIO`. For example, if `OS_LOWEST_PRIO` is 10 then the matrix is actually 2x8 (two bytes of 8 bits). Similarly, if `OS_LOWEST_PRIO` is set to 47, the matrix will be 6x8. When `OS_LOWEST_PRIO` is above 63, we use 16-bit wide entries. For example, if you specify `OS_LOWEST_PRIO` to be 100 then the matrix will be 7x16 (7 entries of 16 bits each). You **CANNOT** set `OS_LOWEST_PRIO` to 255 because this value is reserved for `OS_PRIO_SELF`.

New Files

APP_CFG.H

(Added in V2.80)

We now assume the presence of a file called APP_CFG.H which is declared in your application. The purpose of this file is to assign task priorities, stack sizes and other configuration information for your application.

OS_CFG_R.H

(Added in V2.70)

This file is 'reference' file so that you don't have to create this file from scratch. OS_CFG_R.H has been added in V2.70 and is found in the 'Source' directory of the microprocessor independent portion of **µC/OS-II**. It is recommended that you *copy* OS_CFG_R.H to OS_CFG.H of your project directory.

OS_TMR.C

(Added in V2.81, revised in V2.83)

We added a timer manager function in **µC/OS-II**. You can now define any number of timers. The timers can be periodic or one-shots. A user definable function can be called when the timer expires. One such function is definable for each timer in your application.

New Port Files

OS_DBG.C

(Added in V2.62 but renamed from OS_DEBUG.C in V2.70)

OS_DBG_R.C

(Added in V2.70)

This file should be placed in the same directory as OS_CPU_C.C, OS_CPU.H and OS_CPU_A.ASM of the port you are using. OS_DBG.C defines a series of variables that are placed in ROM (code space). These variables are used by some Kernel Aware Debuggers to get information about **µC/OS-II** and its configuration. If you DON'T use a Kernel Aware Debugger that requires this file, you DON'T need to have it. Check you Kernel Aware Debugger documentation. OS_DBG.C used to be called OS_DEBUG.C in V2.62.

OS_DBG_R.C is a 'reference' file so that you don't have to create this file from scratch. OS_DBG_R.C has been added in V2.70 and is found in the 'Source' directory of the microprocessor independent portion of **µC/OS-II**.

Changes

`uCOS_II.H`

(Changed in V2.70, V2.80 and V2.84)

This file now includes `#include` statements to include `APP_CFG.H`, `OS_CPU.H` and `OS_CFG.H`. This allows you to compile **μC/OS-II** without the needs of any other library functions.

Chaned error codes to make them more consistent. Specifically, all error codes start with `OS_ERR_`. The old error codes have been kept for backward compatibility but you should now use and check for the new error codes.

Time delays and Timeouts

(Changed in V2.87)

All time delays and timeouts are now implemented using an unsigned 32-bit variable. This actually simplified `OSTimeDlyHMSM()` and allowed `OSTimeDlyResume()` to work with any delay. Of course, this means that additional storage is needed in the tasks `OS_TCBs` (2 extra bytes) but that should not be a problem with most applications. It turns out that close to 90% of **μC/OS-II** users use 32-bit CPUs.

Names of objects stored as pointers

(Changed in V2.87)

Names of objects were previously stored in RAM inside the different kernel objects. Specifically, RAM storage was allocated in the `OS_TCB`, for example, to store the name of a task. As of V2.87, all such names are now referenced using pointers. This drastically reduces the amount of RAM needed to store ASCII names since names are now typically allocated by the compiler as constant strings and thus placed in ROM. This was done to reduce the amount of RAM needed (a processor typically has more ROM than RAM) and also to lift the limit of the length of a kernel object name.

Stack check usage in number of entries

(Changed in V2.92)

To be consistent with `OSTaskCreate()` and `OSTaskCreateExt()`, `OSTaskStkChk()` reports stack usage in number of stack entries (or elements) as opposed to number of bytes. To convert number of entries in number of bytes, multiply the value by `sizeof(OS_STK)`.

New #define Constants and Macros

`OS_APP_HOOKS_EN` (OS_CFG.H, V2.85)

This constant specifies whether **μC/OS-II**'s hook functions will call application defined hooks. Specifically, when set to 1 ...

| The μC/OS-II hook ... | Calls the Application-define hook ... |
|---------------------------------|---------------------------------------|
| <code>OSTaskCreateHook()</code> | <code>App_TaskCreateHook()</code> |
| <code>OSTaskDelHook()</code> | <code>App_TaskDelHook()</code> |
| <code>OSTaskIdleHook()</code> | <code>App_TaskIdleHook()</code> |
| <code>OSTaskStatHook()</code> | <code>App_TaskStatHook()</code> |
| <code>OSTaskSwHook()</code> | <code>App_TaskSwHook()</code> |
| <code>OSTCBInitHook()</code> | <code>App_TCBInitHook()</code> |
| <code>OSTimeTickHook()</code> | <code>App_TimeTickHook()</code> |

`OS_ARG_CHK_EN` (OS_CFG.H, V2.04)

This constant is used to specify whether argument checking will be performed at the beginning of most of **μC/OS-II** services. You should always choose to turn this feature on (when set to 1) unless you need to get the best performance possible out of **μC/OS-II** or, you need to reduce code size.

`OS_CRITICAL_METHOD #3` (OS_CPU.H, V2.04)

This constant specifies the method used to disable and enable interrupts during critical sections of code. Prior to V2.04, `OS_CRITICAL_METHOD` could be set to either 1 or 2. In V2.04, I added a local variable (i.e. `cpu_sr`) in most function calls to save the processor status register which generally holds the state of the interrupt disable flag(s). You would then declare the two critical section macros as follows:

```
#define OS_ENTER_CRITICAL() (cpu_sr = OS_CPUSR_Save())
#define OS_EXIT_CRITICAL() (OS_CPU_SR_Restore(cpu_sr))
```

Note that the functions `OS_CPU_SR_Save()` and `OS_CPU_SR_Restore()` would be written in assembly language and would typically be found in `OS_CPU_A.ASM` (or equivalent).

`OS_DEBUG_EN` (OS_CFG.H, V2.60)

This constant is used to enable ROM constants used for debugging using a kernel aware debugger. The constants are found in `OS_CORE.C`.

`OS_EVENT_MULTII_EN` (OS_CFG.H, V2.86)

This constant determines whether the code to support pending on multiple events will be enabled (1) or not (0). This constant thus enables code for the function `OSEventPendMulti()`. This #define was added in V2.86.

`OS_EVENT_NAME_EN` (OS_CFG.H, V2.60 and changed in V2.87)

This constant determines whether names can be assigned to a semaphore, a mutex, a mailbox or a message queue. If `OS_EVENT_NAME_EN` is set to 0, this feature is disabled.

`OS_FLAG_EN` (OS_CFG.H, V2.51)

This constant is used to specify whether you will enable (when 1) code generation for the event flags.

`OS_FLAG_NAME_EN` (OS_CFG.H, V2.60 and changed in V2.87)

This constant determines whether names can be assigned to event flag groups. If `OS_FLAG_NAME_EN` is set to 0, this feature is disabled.

`OS_FLAG_WAIT_CLR_EN` (OS_CFG.H, V2.51)

This constant is used to enable code generation (when 1) to allow to wait on cleared event flags.

`OS_MAX_FLAGS` (OS_CFG.H, V2.51)

This constant is used to determine how many event flags your application will support.

`OS_MBOX_PEND_ABORT_EN` (OS_CFG.H, V2.84)

This constant is used to determine whether you will enable (when 1) code generation for the `OSMboxPendAbort()` function.

`OS_MEM_NAME_EN` (OS_CFG.H, V2.60 and changed in V2.87)

This constant determines whether names can be assigned to memory partitions. If `OS_MEM_NAME_EN` is set to 0, this feature is disabled and no RAM is used in the `OS_MEM` for the memory partition.

`OS_MUTEX_EN` (OS_CFG.H, V2.04)

This constant is used to specify whether you will enable (when 1) code generation for mutual exclusion semaphores.

`OS_Q_PEND_ABORT_EN` (OS_CFG.H, V2.84)

This constant is used to determine whether you will enable (when 1) code generation for the `OSQPendAbort()` function.

`OS_SEM_PEND_ABORT_EN` (OS_CFG.H, V2.84)

This constant is used to determine whether you will enable (when 1) code generation for the `OSSemPendAbort()` function.

OS_TASK_NAME_EN (OS_CFG.H, V2.60 and changed in V2.87)

This constant determines whether names can be assigned to tasks. If OS_TASK_NAME_EN is set to 0, this feature is disabled and no RAM is used in the OS_TCB for the task name.

OS_TASK_PROFILE_EN (OS_CFG.H, V2.60)

This constant allows variables to be allocated in each task's OS_TCB that hold performance data about each task. Specifically, if OS_TASK_PROFILE_EN is set to 1, each task will have a variable to keep track of the number of context switches, the task execution time, the number of bytes used by the task and more.

OS_TASK_STAT_STK_CHK_EN (OS_CFG.H, V2.60)

This constant allows the statistic task to determine the actual stack usage of each active task. If OS_TASK_STAT_EN is set to 0 (the statistic task is not enabled), you can call OS_TaskStatStkChk() yourself from one of your tasks. . If OS_TASK_STAT_EN is set to 1, stack sizes will be determined every second.

OS_TASK_SW_HOOK_EN (OS_CFG.H, V2.60)

Normally, **µC/OS-II** requires that you have a context switch hook function called OSTaskSwHook(). When set to 0, this constant allows you to omit OSTaskSwHook() from your code. This configuration constant was added to reduce the amount of overhead during a context switch in applications that doesn't require the context switch hook. Of course, you will also need to remove the calls to OSTaskSwHook() from OSTaskStartHighRdy(), OSCtxSw() and OSIntCtxSw() in OS_CPU_A.ASM.

OS_TASK_TMR_STK_SIZE (OS_CFG.H, V2.81)

This #define determines the stack size (in number of stack-size elements, i.e. OS_STK) of the timer task. The size of the timer task's stack greatly depends on the processor architecture and the functions that are called when timers expire. Note that if you set OS_TMR_EN to 0 in OS_CFG.H then the value you set for OS_TASK_TMR_STK_SIZE is irrelevant because the timer functionality would be disabled.

OS_TICK_STEP_EN (OS_CFG.H, V2.60)

µC/OS-View can now 'halt' **µC/OS-II**'s tick processing and allow you to issue 'step' commands from **µC/OS-View**. In other words, **µC/OS-View** can prevent **µC/OS-II** from calling OSTimeTick() so that timeouts and time delays are no longer processed. However, though a keystroke from **µC/OS-View**, you can execute a single tick at a time. If OS_TIME_TICK_HOOK_EN (see below) is set to 1, OSTimeTickHook() is still executed at the regular tick rate in case you have time critical items to take care of in your application.

OS_TIME_TICK_HOOK_EN (OS_CFG.H, V2.60)

Normally, **µC/OS-II** requires the presence of a function called OSTimeTickHook() which is called at the very beginning of the tick ISR. When set to 0, this constant allows you to omit OSTimeTickHook() from your code. This configuration constant was added to reduce the amount of overhead during a tick ISR in applications that doesn't require this hook.

OS_TMR_EN (OS_CFG.H, V2.81)

This #define enables (when set to 1) or disables (when set to 0) the timer management code.

OS_TMR_CFG_MAX (OS_CFG.H, V2.81)

This #define determines the maximum number of timers that can exist in the application. If OS_TMR_EN is set to 1, you should declare **AT LEAST** two (2) timers.

OS_TMR_CFG_NAME_EN (OS_CFG.H, V2.81 and changed in V2.87)

This #define determines whether names can be assigned to timers.

OS_TMR_CFG_WHEEL_SIZE (OS_CFG.H, V2.81)

This #define determines the number of entries in the timer wheel. This value should be a number between 2 and 1024. Timer management overhead is somewhat determined by the size of the wheel. A large number of entries might reduce the overhead for timer management but would require more RAM. Each entry requires a pointer and a count (16-bit value). We recommend a number that is NOT a multiple of the tick rate. If your application has many timers then it's recommended that you have a high value. As a starting value, you could use OS_TMR_CFG_MAX / 4.

OS_TMR_CFG_TICKS_PER_SEC (OS_CFG.H, V2.81)

This #define determines the rate at which timers will be updated. You would typically set to a fraction of the tick rate (i.e. OS_TICKS_PER_SEC). We recommend that you set OS_TMR_CFG_TICKS_PER_SEC to 10 (i.e. 10 Hz).

The following table summarizes some of the new #define constants in OS_CFG.H which were all added in since V2.00.

| #define name in OS_CFG.H | ... to enable the function(s): |
|--------------------------|--|
| OS_APP_HOOKS_EN | App_TaskCreateHook() App_TaskDelHook() App_TaskIdleHook() App_TaskStatHook() App_TaskSwHook() App_TCBInitHook() App_TimeTickHook() |
| OS_DEBUG_EN | Enable debug constants in OS_CORE.C. If you are using a kernel aware debugger, you should enable this feature. |
| | |
| OS_EVENT_NAME_EN | OSEventNameGet() OSEventNameSet() And, to allow naming semaphores, mutexes, mailboxes and message queues. |
| OS_EVENT_MULTI_EN | OSEventPendMulti() |
| | |
| OS_FLAG_ACCEPT_EN | OSFlagAccept() |
| OS_FLAG_DEL_EN | OSFlagDel() |
| OS_FLAG_NAME_EN | OSFlagNameGet() OSFlagNameSet() And, to allow naming event flag groups. |
| OS_FLAG_QUERY_EN | OSFlagQuery() |
| | |
| OS_MBOX_ACCEPT_EN | OSMboxAccept() |
| OS_MBOX_DEL_EN | OSMboxDel() |
| OS_MBOX_PEND_ABORT_EN | OSMboxPendAbort() |
| OS_MBOX_POST_EN | OSMboxPost() |
| OS_MBOX_POST_OPT_EN | OSMboxPostOpt() |
| OS_MBOX_QUERY_EN | OSMboxQuery() |
| | |
| OS_MEM_NAME_EN | OSMemNameGet() OSMemNameSet() |
| OS_MEM_QUERY_EN | OSMemQuery() |
| | |
| OS_MUTEX_ACCEPT_EN | OSMutexAccept() |
| OS_MUTEX_DEL_EN | OSMutexDel() |
| OS_MUTEX_QUERY_EN | OSMutexQuery() |

| | |
|--------------------------|--|
| OS_Q_ACCEPT_EN | OSQAccept() |
| OS_Q_DEL_EN | OSQDel() |
| OS_Q_FLUSH_EN | OSQFlush() |
| OS_Q_PEND_ABORT_EN | OSQPendAbort() |
| OS_Q_POST_EN | OSQPost() |
| OS_Q_POST_FRONT_EN | OSQPostFront() |
| OS_Q_POST_OPT_EN | OSQPostOpt() |
| OS_Q_QUERY_EN | OSQQuery() |
| | |
| OS_SEM_ACCEPT_EN | OSSemAccept() |
| OS_SEM_DEL_EN | OSSemDel() |
| OS_SEM_PEND_ABORT_EN | OSSemPendAbort() |
| OS_SEM_QUERY_EN | OSSemQuery() |
| OS_SEM_SET_EN | OSSemSet() |
| | |
| OS_TASK_NAME_EN | OSTaskNameGet() OSTaskNameSet() And, to allow naming tasks. |
| OS_TASK_PROFILE_EN | To allocate variables in OS_TCB for performance monitoring of each task at run-time. |
| OS_TASK_QUERY_EN | OSTaskQuery() |
| OS_TASK_STAT_STK_CHK_EN | OS_TaskStatStkChk() |
| OS_TASK_SW_HOOK_EN | OSTaskSwHook() |
| OS_TASK_TMR_STK_SIZE | Size in OS_STK elements of the Timer Management task. |
| | |
| OS_TICK_STEP_EN | To support the stepping feature of μC/OS-View . |
| | |
| OS_TIME_DLY_HMSM_EN | OSTimeDlyHMSM() |
| OS_TIME_DLY_RESUME_EN | OSTimeDlyResume() |
| OS_TIME_GET_SET_EN | OSTimeGet() and OSTimeSet() |
| OS_TIME_TICK_HOOK_EN | OSTimeTickHook() |
| | |
| OS_TMR_EN | Enables (1) or Disables (0) timer management functions. |
| OS_TMR_CFG_MAX | Determines the maximum number of timers in your application. |
| OS_TMR_CFG_NAME_EN | Determines whether names can be assigned to timers. |
| OS_TMR_CFG_WHEEL_SIZE | Determines the size of the timer wheel (in number of entries). |
| OS_TMR_CFG_TICKS_PER_SEC | Rate at which timers will be updated (Hz) |
| | |
| OS_SCHED_LOCK_EN | OSSchedLock() and OSSchedUnlock() |

New Data Types

`OS_CPU_SR` (OS_CPU.H, V2.04)

This data type is used to specify the size of the CPU status register which is used in conjunction with `OS_CRITICAL_METHOD #3` (see above). For example, if the CPU status register is 16-bit wide then you would typedef accordingly.

`OS_FLAGS` (uCOS_II.H, V2.51)

This data type determines how many bits an event flag group will have. You can thus typedef this data type to either `INT8U`, `INT16U` or `INT32U` to give event flags either 8, 16 or 32 bits, respectively.

`OS_TMR` (uCOS_II.H, V2.81)

This data type is a timer object which contains information about a specific timer that you started (see `OS_TMR.C`).

New Hook Functions

`void OSInitHookBegin(void)` (OS_CPU.C, V2.04)

This function is called at the very beginning of `OSInit()` to allow for port specific initialization BEFORE **μC/OS-II** gets initialized.

`void OSInitHookEnd(void)` (OS_CPU.C, V2.04)

This function is called at the end of `OSInit()` to allow for port specific initialization AFTER **μC/OS-II** gets initialized.

`void OSTCBInitHook(OS_TCB *ptcb)` (OS_CPU.C, V2.04)

This function is called by `OSTCBInit()` during initialization of the TCB assigned to a newly created task. It allows port specific initialization of the TCB.

`void OSTaskIdleHook(void)` (OS_CPU.C, V2.51)

This function is called by `OSTaskIdle()`. This allows you to STOP the CPU and thus reduce power consumption while there is nothing to do.

New Functions

The following table provides a list of all the new functions (i.e. services) that YOUR application can call. The list also includes functions that used to exist but, if these are in this list, it's because their API may have changed.

Refer to the *Reference Manual* of the current release for a description of these functions.

| Function Name | File | Enabled By ... |
|-------------------------|------------|-------------------------------------|
| OSEventNameGet() | OS CORE.C | OS EVENT NAME EN |
| OSEventNameSet() | OS CORE.C | OS EVENT NAME EN |
| OSEventPendMulti() | OS CORE.C | OS EVENT MULTI EN |
| OSFlagAccept() | OS FLAG.C | OS FLAG EN && OS FLAG ACCEPT EN |
| OSFlagCreate() | OS FLAG.C | OS FLAG EN |
| OSFlagDel() | OS FLAG.C | OS FLAG EN && OS FLAG DEL EN |
| OSFlagNameGet() | OS FLAG.C | OS FLAG NAME EN |
| OSFlagNameSet() | OS FLAG.C | OS FLAG NAME EN |
| OSFlagPend() | OS FLAG.C | OS FLAG EN |
| OSFlagPendGetFlagsRdy() | OS FLAG.C | OS FLAG EN |
| OSFlagPost() | OS FLAG.C | OS FLAG EN |
| OSFlagQuery() | OS FLAG.C | OS FLAG EN |
| OSMboxDel() | OS MBOX.C | OS MBOX EN && OS MBOX DEL EN |
| OSMboxPendAbort() | OS MBOX.C | OS MBOX EN && OS MBOX PEND ABORT EN |
| OSMboxPostOpt() | OS MBOX.C | OS MBOX EN && OS MBOX POST OPT EN |
| OSMutexAccept() | OS MUTEX.C | OS MUTEX EN && OS MUTEX ACCEPT EN |
| OSMutexCreate() | OS MUTEX.C | OS MUTEX EN |
| OSMutexDel() | OS MUTEX.C | OS MUTEX EN && OS MUTEX DEL EN |
| OSMutexPend() | OS MUTEX.C | OS MUTEX EN |
| OSMutexPost() | OS MUTEX.C | OS MUTEX EN |
| OSMutexQuery() | OS MUTEX.C | OS MUTEX EN && OS MUTEX QUERY EN |
| OSQAccept() | OS Q.C | OS Q EN && OS Q ACCEPT EN |
| OSQDel() | OS Q.C | OS Q EN && OS Q DEL EN |
| OSQFlush() | OS Q.C | OS Q EN && OS Q FLUSH EN |
| OSQPend() | OS Q.C | OS Q EN |
| OSQPendAbort() | OS Q.C | OS Q EN && OS Q PEND ABORT EN |
| OSQPost() | OS Q.C | OS Q EN |
| OSQPostFront() | OS Q.C | OS Q EN && OS Q POST FRONT EN |
| OSQPostOpt() | OS Q.C | OS Q EN && OS Q POST OPT EN |
| OSSafetyCriticalStart() | OS CORE.C | OS SAFETY CRITICAL IEC61508 |
| OSSemDel() | OS SEM.C | OS SEM EN && OS SEM DEL EN |
| OSSemPendAbort() | OS SEM.C | OS SEM EN && OS SEM PEND ABORT EN |
| OSSemSet() | OS SEM.C | OS SEM EN && OS SEM SET EN |
| OSTaskNameGet() | OS TASK.C | OS TASK NAME EN |
| OSTaskNameSet() | OS TASK.C | OS TASK NAME EN |
| OSTmrGetName() | OS TMR.C | OS TMR EN |
| OSTmrGetRemain() | OS TMR.C | OS TMR EN |
| OSTmrStart() | OS TMR.C | OS TMR EN |
| OSTmrStop() | OS TMR.C | OS TMR EN |
| OSTmrSignal() | OS TMR.C | OS TMR EN |

References

μC/OS-II, The Real-Time Kernel, 2nd Edition

Jean J. Labrosse
CMP Books, 2002
ISBN 1-57820-103-9

Contacts

Micrium

949 Crestview Circle
Weston, FL 33327
954-217-2036
954-217-2037 (FAX)
e-mail: Jean.Labrosse@Micrium.com
WEB: www.Micrium.com

CMP Books, Inc.

1601 W. 23rd St., Suite 200
Lawrence, KS 66046-9950
(785) 841-1631
(785) 841-2624 (FAX)
WEB: <http://www.cmpbooks.com>
e-mail: rdorders@cmpbooks.com